

SE274 Data Structure


Lecture 6: Maps, Hash Tables, Skip Lists – Part 2

Apr 08, 2020

Instructor: Sunjun Kim

Information&Communication Engineering, DGIST

Sets



Definitions

- A **set** is an unordered collection of elements, without duplicates that typically supports efficient membership tests.
 - Elements of a set are like keys of a map, but without any auxiliary values.
- A **multiset** (also known as a **bag**) is a set-like container that allows duplicates.
- A **multimap** is similar to a traditional map, in that it associates values with keys; however, in a multimap the same key can be mapped to multiple values.
 - For example, the index of a book maps a given term to one or more locations at which the term occurs.

Set ADT

`S.add(e)`: Add element `e` to the set. This has no effect if the set already contains `e`.

`S.discard(e)`: Remove element `e` from the set, if present. This has no effect if the set does not contain `e`.

`e in S`: Return `True` if the set contains element `e`. In Python, this is implemented with the special `__contains__` method.

`len(S)`: Return the number of elements in set `S`. In Python, this is implemented with the special method `__len__`.

`iter(S)`: Generate an iteration of all elements of the set. In Python, this is implemented with the special method `__iter__`.

`S.remove(e)`: Remove element `e` from the set. If the set does not contain `e`, raise a `KeyError`.

`S.pop()`: Remove and return an arbitrary element from the set. If the set is empty, raise a `KeyError`.

`S.clear()`: Remove all elements from the set.

Boolean Set Operations

$S == T$: Return True if sets S and T have identical contents.

$S != T$: Return True if sets S and T are not equivalent.

$S \leq T$: Return True if set S is a subset of set T.

$S < T$: Return True if set S is a *proper* subset of set T.

$S \geq T$: Return True if set S is a superset of set T.

$S > T$: Return True if set S is a *proper* superset of set T.

$S.isdisjoint(T)$: Return True if sets S and T have no common elements.

Set Update Operations

$S \mid T$: Return a new set representing the union of sets S and T .

$S \mid= T$: Update set S to be the union of S and set T .

$S \& T$: Return a new set representing the intersection of sets S and T .

$S \&= T$: Update set S to be the intersection of S and set T .

$S \wedge T$: Return a new set representing the symmetric difference of sets S and T , that is, a set of elements that are in precisely one of S or T .

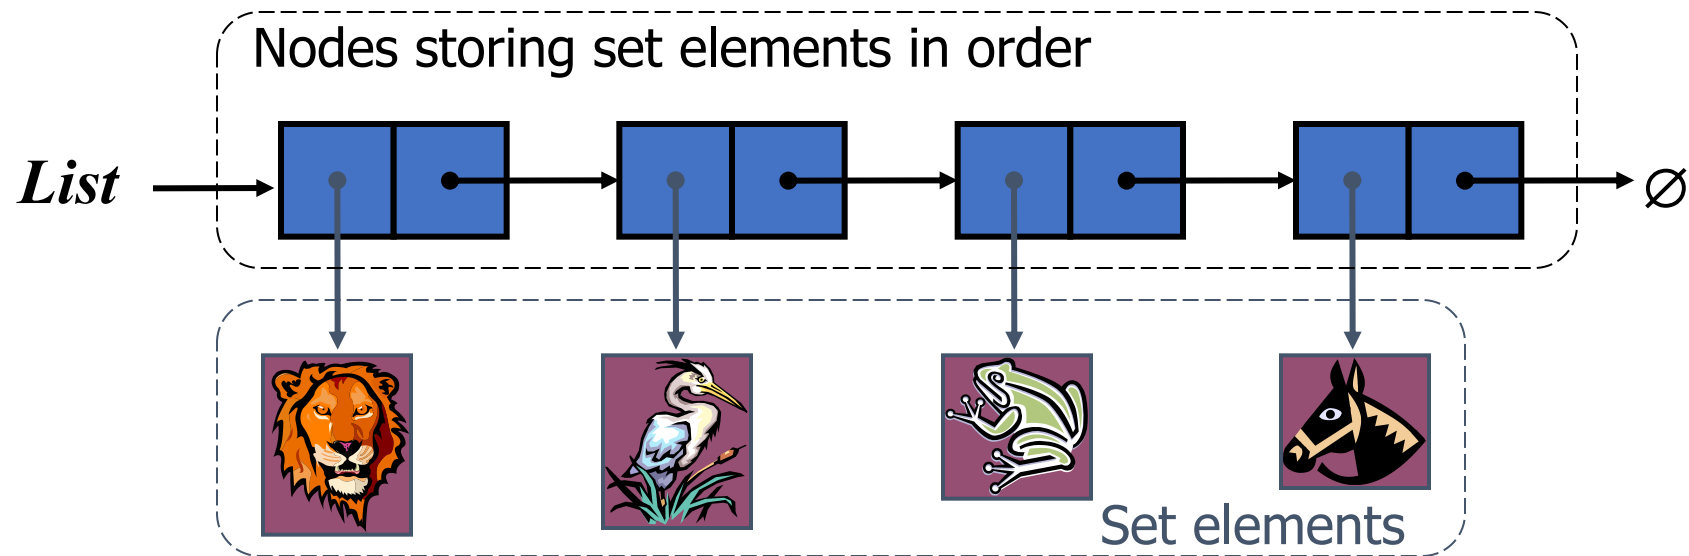
$S \wedge= T$: Update set S to become the symmetric difference of itself and set T .

$S - T$: Return a new set containing elements in S but not T .

$S -= T$: Update set S to remove all common elements with set T .

Storing a Set in a List

- We can implement a set with a list
- Elements are stored sorted according to some canonical ordering
- The space used is $O(n)$



Generic Merging

- Generalized merge of two sorted lists A and B
- Template method `genericMerge`
- Auxiliary methods
 - `alsLess`
 - `blsLess`
 - `bothAreEqual`
- Runs in $O(n_A + n_B)$ time provided the auxiliary methods run in $O(1)$ time

Algorithm *genericMerge*(A, B)

$S \leftarrow$ empty sequence

while $\neg A.isEmpty() \wedge \neg B.isEmpty()$

$a \leftarrow A.first().element(); b \leftarrow B.first().element()$

if $a < b$

$aIsLess(a, S); A.remove(A.first())$

else if $b < a$

$bIsLess(b, S); B.remove(B.first())$

else $\{ b = a \}$

$bothAreEqual(a, b, S)$

$A.remove(A.first()); B.remove(B.first())$

while $\neg A.isEmpty()$

$aIsLess(a, S); A.remove(A.first())$

while $\neg B.isEmpty()$

$bIsLess(b, S); B.remove(B.first())$


return S

Using Generic Merge for Set Operations




- Any of the set operations can be implemented using a generic merge
- For example:
 - For **intersection**: only copy elements that are duplicated in both list
 - For **union**: copy every element from both lists except for the duplicates
- All methods run in linear time

Skip Lists




What is a Skip List

- A **skip list** for a set S of distinct (key, element) items is a series of lists S_0, S_1, \dots, S_h such that
 - Each list S_i contains the special keys $+\infty$ and $-\infty$
 - List S_0 contains the keys of S in nondecreasing order
 - Each list is a subsequence of the previous one, i.e.,
 $S_0 \supseteq S_1 \supseteq \dots \supseteq S_h$
 - List S_h contains only the two special keys
- We show how to use a skip list to implement the dictionary ADT




Actions in a skip list

- `next(p)` : Return the position following **p** on the same level.
 - `prev(p)` : Return the position preceding **p** on the same level.
 - `below(p)` : Return the position below **p** in the same tower.
 - `above(p)` : Return the position above **p** in the same tower.
- * The actions return **None** if the position requested does not exist



Search

- We search for a key x in a skip list as follows:
 - We start at the first position of the top list
 - At the current position p , we compare x with $y \leftarrow \text{key}(\text{next}(p))$
 - $x = y$: we return $\text{element}(\text{next}(p))$
 - $x > y$: we “scan forward”
 - $x < y$: we “drop down”
 - If we try to drop down past the bottom list, we return *null*
- Example: search for 78



Search Pseudo-code

Algorithm SkipSearch(k):

Input: A search key k

Output: Position p in the bottom list S_0 with the largest key such that $\text{key}(p) \leq k$

$p = \text{start}$

{begin at start position}

while $\text{below}(p) \neq \text{None}$ **do**

$p = \text{below}(p)$

{drop down}

while $k \geq \text{key}(\text{next}(p))$ **do**

$p = \text{next}(p)$

{scan forward}

return p .

Code Fragment 10.12: Algorithm to search a skip list S for key k .


Randomized Algorithms

- A randomized algorithm performs coin tosses (i.e., uses random bits) to control its execution
- It contains statements of the type

```
b ← random()
if b = 0
  do A ...
else { b = 1 }
  do B ...
```
- Its running time depends on the outcomes of the coin tosses
- We analyze the expected running time of a randomized algorithm under the following assumptions
 - the coins are unbiased, and
 - the coin tosses are independent
- The worst-case running time of a randomized algorithm is often large but has very low probability (e.g., it occurs when all the coin tosses give “heads”)
- We use a randomized algorithm to insert items into a skip list

Insertion

- To insert an entry (x, o) into a skip list, we use a randomized algorithm:
 - We repeatedly toss a coin until we get tails, and we denote with i the number of times the coin came up heads
 - If $i \geq h$, we add to the skip list new lists S_{h+1}, \dots, S_{i+1} , each containing only the two special keys
 - We search for x in the skip list and find the positions p_0, p_1, \dots, p_i of the items with largest key less than x in each list S_0, S_1, \dots, S_i
 - For $j \leftarrow 0, \dots, i$, we insert item (x, o) into list S_j after position p_j
- Example: insert key 15, with $i = 2$



Insertion Pseudo-code

Algorithm SkipInsert(k, v):

Input: Key k and value v

Output: Topmost position of the item inserted in the skip list

$p = \text{SkipSearch}(k)$

$q = \text{None}$

{ q will represent top node in new item's tower}

$i = -1$

repeat

$i = i + 1$

if $i \geq h$ **then**

$h = h + 1$

{add a new level to the skip list}

$t = \text{next}(s)$

$s = \text{insertAfterAbove}(\text{None}, s, (-\infty, \text{None}))$

{grow leftmost tower}

$\text{insertAfterAbove}(s, t, (+\infty, \text{None}))$

{grow rightmost tower}

while $\text{above}(p)$ is **None** **do**

$p = \text{prev}(p)$

{scan backward}

$p = \text{above}(p)$

{jump up to higher level}

$q = \text{insertAfterAbove}(p, q, (k, v))$

{increase height of new item's tower}

until $\text{coinFlip}() == \text{tails}$

$n = n + 1$

return q

$r = \text{insertAfterAbove}(p, q, (k, v))$
- Inserts a position storing the item (k, v) after position p (on the same level as p) and above position q , returning the new position r .


n : # of entries

h : the height

s : the start node


Deletion

- To remove an entry with key x from a skip list, we proceed as follows:
 - We search for x in the skip list and find the positions p_0, p_1, \dots, p_i of the items with key x , where position p_j is in list S_j
 - We remove positions p_0, p_1, \dots, p_i from the lists S_0, S_1, \dots, S_i
 - We remove all but one list containing only the two special keys
- Example: remove key 34



Implementation

- We can implement a skip list with quad-nodes
- A quad-node stores:
 - entry
 - link to the node prev
 - link to the node next
 - link to the node below
 - link to the node above
- Also, we define special keys PLUS_INF and MINUS_INF, and we modify the key comparator to handle them



Space Usage

- The space used by a skip list depends on the random bits used by each invocation of the insertion algorithm
- We use the following two basic probabilistic facts:
 - Fact 1: The probability of getting i consecutive heads when flipping a coin is $1/2^i$
 - Fact 2: If each of n entries is present in a set with probability p , the expected size of the set is np
- Consider a skip list with n entries
 - By Fact 1, we insert an entry in list S_i with probability $1/2^i$
 - By Fact 2, the expected size of list S_i is $n/2^i$
- The expected number of nodes used by the skip list is

$$\sum_{i=0}^h \frac{n}{2^i} = n \sum_{i=0}^h \frac{1}{2^i} < 2n$$

- ◆ Thus, the expected space usage of a skip list with n items is $O(n)$

Height

- The running time of the search and insertion algorithms is affected by the height h of the skip list
- We show that with high probability, a skip list with n items has height $O(\log n)$
- We use the following probabilistic fact:
 - Fact 3: If each of n events has probability p , the probability that at least one event occurs is at most np
- Consider a skip list with n entries
 - By Fact 1, we insert an entry in list S_i with probability $1/2^i$
 - By Fact 3, the probability that list S_i has at least one item is at most $n/2^i$
- By picking $i = 3\log n$, we have that the probability that $S_{3\log n}$ has at least one entry is at most
$$n/2^{3\log n} = n/n^3 = 1/n^2$$
- Thus a skip list with n entries has height at most $3\log n$ with probability at least $1 - 1/n^2$

Search and Update Times

- The search time in a skip list is proportional to
 - the number of drop-down steps, plus
 - the number of scan-forward steps
- The drop-down steps are bounded by the height of the skip list and thus are $O(\log n)$ with high probability
- To analyze the scan-forward steps, we use yet another probabilistic fact:
 - Fact 4: The expected number of coin tosses required in order to get tails is 2
- When we scan forward in a list, the destination key does not belong to a higher list
 - A scan-forward step is associated with a former coin toss that gave tails
- By Fact 4, in each list the expected number of scan-forward steps is 2
- Thus, the expected number of scan-forward steps is $O(\log n)$
- We conclude that a search in a skip list takes $O(\log n)$ expected time
- The analysis of insertion and deletion gives similar results

Summary

- A skip list is a data structure for dictionaries that uses a randomized insertion algorithm
- In a skip list with n entries
 - The expected space used is $O(n)$
 - The expected search, insertion and deletion time is $O(\log n)$
- Using a more complex probabilistic analysis, one can show that these performance bounds also hold with high probability
- Skip lists are fast and simple to implement in practice