

SE274 Data Structure

Lecture 7: Sorting and Selection (textbook: Chapter 12)

Apr 06, 2020

Instructor: Sunjun Kim

Information&Communication Engineering, DGIST

Recap: Insertion sort, Selection Sort

- **Selection sort:** Priority Queue
Implementation with an unsorted list



- Performance:
 - **add** takes $O(1)$ time since we can insert the item at the beginning or end of the sequence
 - **Remove_min** and **min** take $O(n)$ time since we have to traverse the entire sequence to find the smallest key


- **Selection sort:** Priority Queue
Implementation with a sorted list




- Performance:
 - **add** takes $O(n)$ time since we have to find the place where to insert the item
 - **remove_min** and **min** take $O(1)$ time, since the smallest key is at the beginning

Recap: heapsort

- **Add** all the items to a heap.
- Repeat **Remove_min** until the heap is empty.
- **Add** takes $O(\log n)$ time, **remove_min** take $O(\log n)$ time.
- Heapsort takes $O(n \log n)$



Merge Sort



Divide-and-Conquer

- Divide-and conquer is a general algorithm design paradigm:
 - Divide: divide the input data S in two disjoint subsets S_1 and S_2
 - Conquer: recursively solve the subproblems associated with S_1 and S_2
 - Combine: combine the solutions for S_1 and S_2 into a solution for S
- The base case for the recursion are subproblems of size 0 or 1
- Merge-sort is a sorting algorithm based on the divide-and-conquer paradigm
- Like heap-sort
 - It has $O(n \log n)$ running time
- Unlike heap-sort
 - It does not use an auxiliary priority queue
 - It accesses data in a sequential manner (suitable to sort data on a disk)

Merge-Sort

- Merge-sort on an input sequence S with n elements consists of three steps:
 - Divide: partition S into two sequences S_1 and S_2 of about $n/2$ elements each
 - Conquer: recursively sort S_1 and S_2
 - Combine: merge S_1 and S_2 into a unique sorted sequence

Algorithm *mergeSort*(S)

Input sequence S with n elements

Output sequence S sorted according to C

if $S.size() > 1$

$(S_1, S_2) \leftarrow partition(S, n/2)$

mergeSort(S_1)

mergeSort(S_2)

$S \leftarrow merge(S_1, S_2)$

Merging Two Sorted Sequences

- The conquer step of merge-sort consists of merging two sorted sequences A and B into a sorted sequence S containing the union of the elements of A and B
- Merging two sorted sequences, each with $n/2$ elements and implemented by means of a doubly linked list, takes $O(n)$ time

Algorithm *merge*(A, B)

Input sequences A and B with $n/2$ elements each

Output sorted sequence of $A \cup B$

$S \leftarrow$ empty sequence

while $\neg A.isEmpty() \wedge \neg B.isEmpty()$
 if $A.first().element() < B.first().element()$
 $S.addLast(A.remove(A.first()))$

else

$S.addLast(B.remove(B.first()))$

while $\neg A.isEmpty()$
 $S.addLast(A.remove(A.first()))$

while $\neg B.isEmpty()$
 $S.addLast(B.remove(B.first()))$

return S


Recap: Generic Merging – set union

- Generalized merge of two sorted lists A and B
- Template method `genericMerge`
- Auxiliary methods
 - `alsLess` => add a
 - `blsLess` => add b
 - `bothAreEqual` => (add b)
- Runs in $O(n_A + n_B)$ time provided the auxiliary methods run in $O(1)$ time

```
Algorithm genericMerge( $A, B$ )
   $S \leftarrow$  empty sequence
  while  $\neg A.isEmpty() \wedge \neg B.isEmpty()$ 
     $a \leftarrow A.first().element(); b \leftarrow B.first().element()$ 
    if  $a < b$ 
       $aIsLess(a, S); A.remove(A.first())$ 
    else if  $b < a$ 
       $bIsLess(b, S); B.remove(B.first())$ 
    else {  $b = a$  }
       $bothAreEqual(a, b, S)$ 
       $A.remove(A.first()); B.remove(B.first())$ 
  while  $\neg A.isEmpty()$ 
     $aIsLess(a, S); A.remove(A.first())$ 
  while  $\neg B.isEmpty()$ 
     $bIsLess(b, S); B.remove(B.first())$ 
  return  $S$ 
```


Python Merge Implementation

```
1 def merge(S1, S2, S):
2     """Merge two sorted Python lists S1 and S2 into properly sized list S."""
3     i = j = 0
4     while i + j < len(S):
5         if j == len(S2) or (i < len(S1) and S1[i] < S2[j]):
6             S[i+j] = S1[i]           # copy ith element of S1 as next item of S
7             i += 1
8         else:
9             S[i+j] = S2[j]           # copy jth element of S2 as next item of S
10            j += 1
```




Python Merge-Sort Implementation

```
1 def merge_sort(S):
2     """Sort the elements of Python list S using the merge-sort algorithm."""
3     n = len(S)
4     if n < 2:
5         return                # list is already sorted
6     # divide
7     mid = n // 2
8     S1 = S[0:mid]              # copy of first half
9     S2 = S[mid:n]              # copy of second half
10    # conquer (with recursion)
11    merge_sort(S1)              # sort copy of first half
12    merge_sort(S2)              # sort copy of second half
13    # merge results
14    merge(S1, S2, S)            # merge sorted halves back into S
```


Merge-Sort Tree

- An execution of merge-sort is depicted by a binary tree
 - each node represents a recursive call of merge-sort and stores
 - unsorted sequence before the execution and its partition
 - sorted sequence at the end of the execution
 - the root is the initial call
 - the leaves are calls on subsequences of size 0 or 1




Execution Example

- Partition




Execution Example (cont.)

- Recursive call, partition




Execution Example (cont.)

- Recursive call, partition




Execution Example (cont.)

- Recursive call, base case




Execution Example (cont.)

- Recursive call, base case




Execution Example (cont.)

- Merge



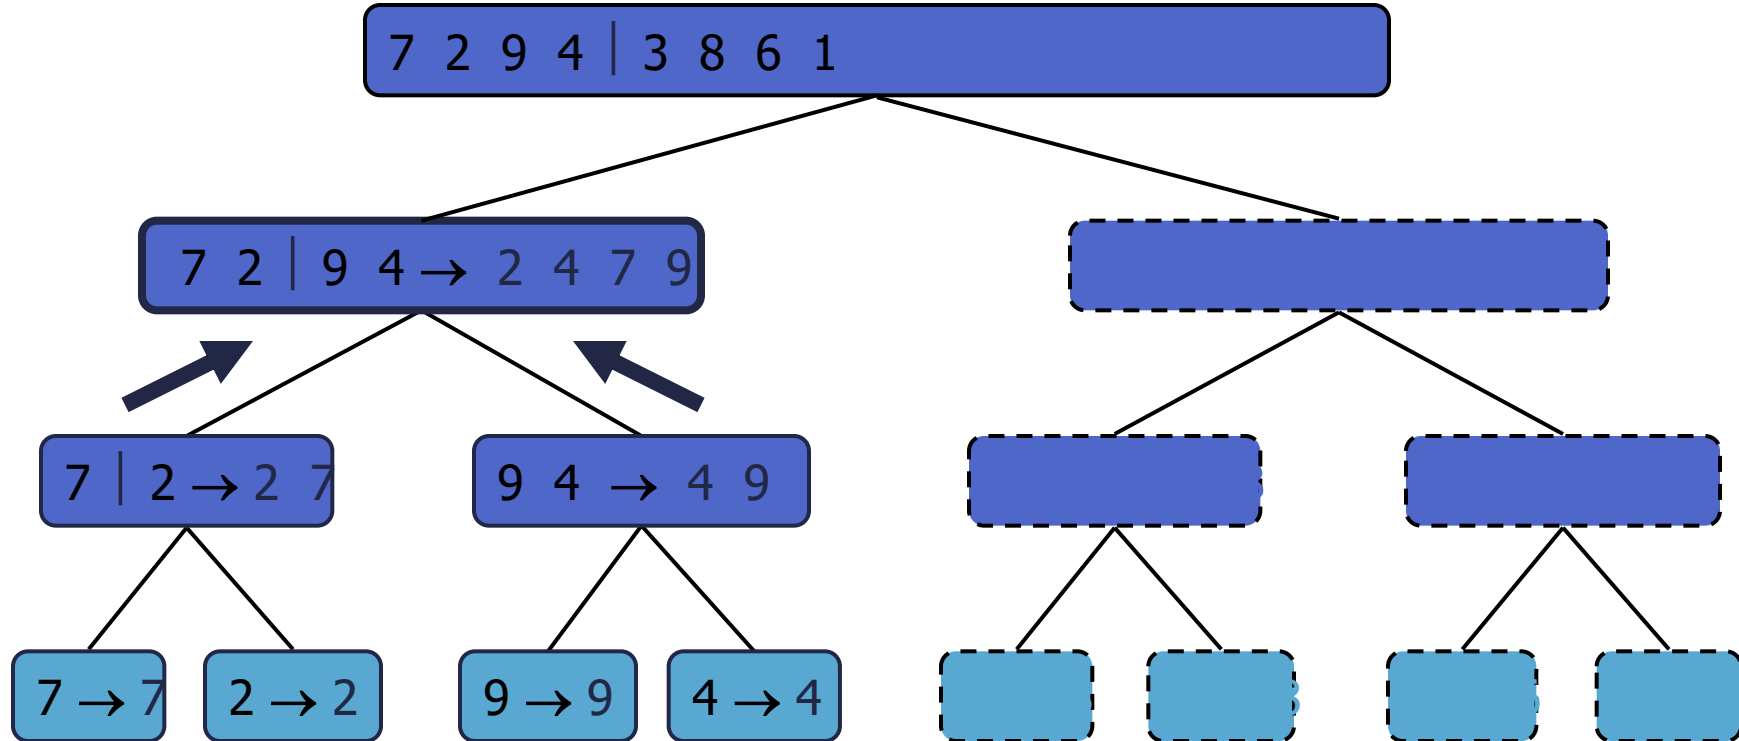
Execution Example (cont.)

- Recursive call, ..., base case, merge



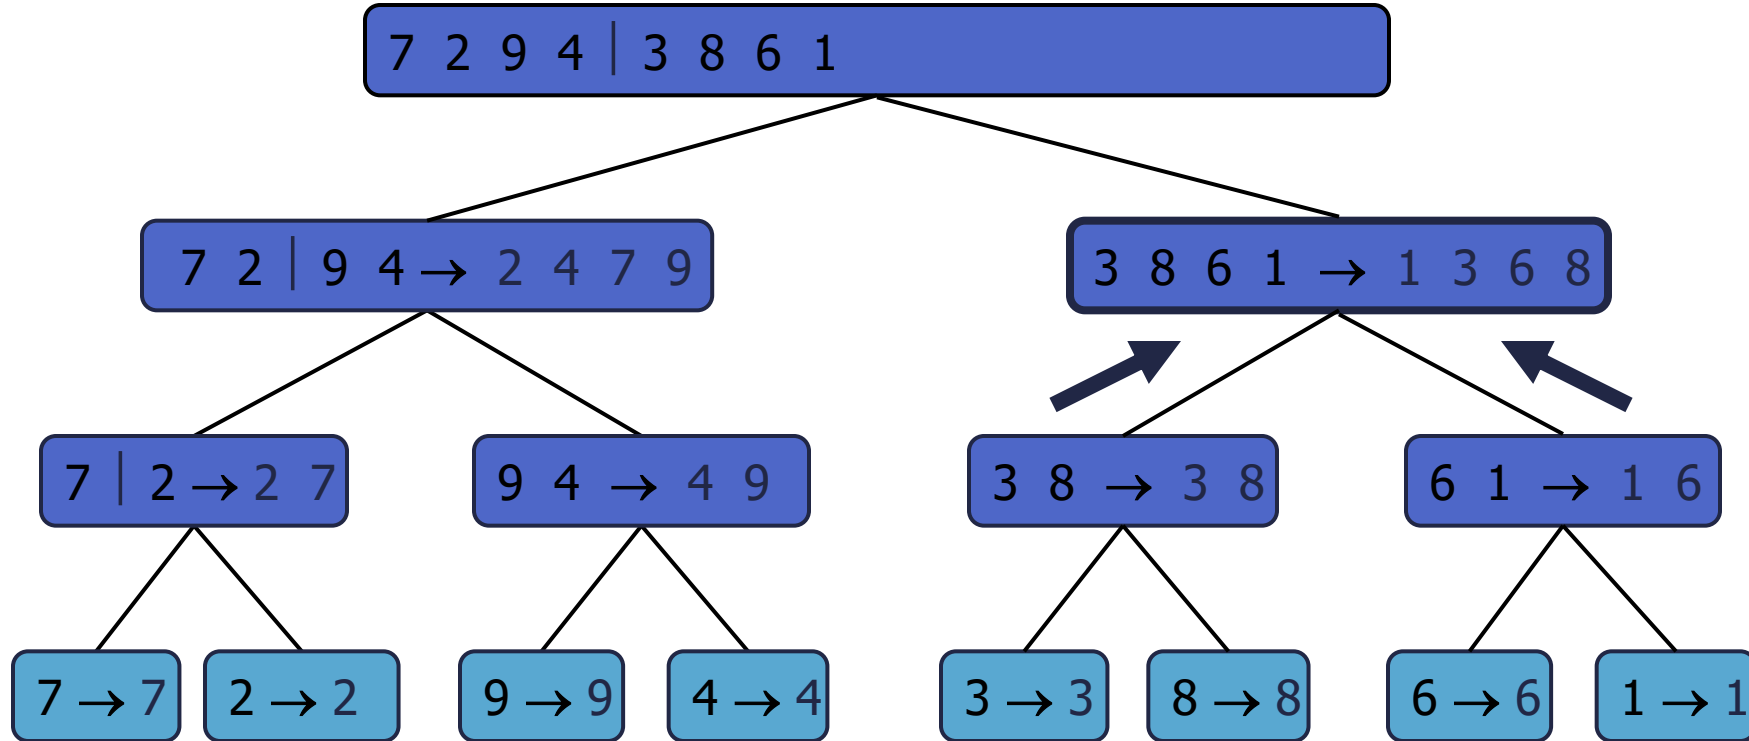
Execution Example (cont.)

- Merge




Execution Example (cont.)

- Recursive call, ..., merge, merge



Execution Example (cont.)

- Merge



Analysis of Merge-Sort

- The height h of the merge-sort tree is $O(\log n)$
 - at each recursive call we divide in half the sequence,
- The overall amount of work done at the nodes of depth i is $O(n)$
 - we partition and merge 2^i sequences of size $n/2^i$
 - we make 2^{i+1} recursive calls
 - When $h = \log n$, $2^{\log n+1} = O(n)$
- Thus, the total running time of merge-sort is $O(n \log n)$


depth #seqs size

0 1 n

1 2 $n/2$

i 2^i $n/2^i$


...



Summary of Sorting Algorithms


| Algorithm | Time | Notes |
|----------------|---------------|--|
| selection-sort | $O(n^2)$ | <ul style="list-style-type: none">▪ slow▪ in-place▪ for small data sets (< 1K) |
| insertion-sort | $O(n^2)$ | <ul style="list-style-type: none">▪ slow▪ in-place▪ for small data sets (< 1K) |
| heap-sort | $O(n \log n)$ | <ul style="list-style-type: none">▪ fast▪ in-place▪ for large data sets (1K — 1M) |
| merge-sort | $O(n \log n)$ | <ul style="list-style-type: none">▪ fast▪ sequential data access▪ for huge data sets (> 1M) |


Quick-Sort



Quick-Sort

- Quick-sort is a randomized sorting algorithm based on the divide-and-conquer paradigm:
 - Divide: pick a random element x (called **pivot**) and partition S into
 - L elements less than x
 - E elements equal x
 - G elements greater than x
 - Conquer: sort L and G
 - Combine: join L , E and G





Partition

- We partition an input sequence as follows:
 - We remove, in turn, each element y from S and
 - We insert y into L , E or G , depending on the result of the comparison with the pivot x
- Each insertion and removal is at the beginning or at the end of a sequence, and hence takes $O(1)$
- Thus, the partition step of quick-sort takes $O(n)$

Algorithm *partition*(S, p)

Input sequence S , position p of pivot

Output subsequences L , E , G of the elements of S less than, equal to, or greater than the pivot, resp.

$L, E, G \leftarrow$ empty sequences

$x \leftarrow S.remove(p)$

while $\neg S.isEmpty()$

$y \leftarrow S.remove(S.first())$

if $y < x$

$L.addLast(y)$

else if $y = x$

$E.addLast(y)$

else $\{ y > x \}$

$G.addLast(y)$


return L, E, G

Python Implementation

```
1 def quick_sort(S):
2     """Sort the elements of queue S using the quick-sort algorithm."""
3     n = len(S)
4     if n < 2:
5         return # list is already sorted
6     # divide
7     p = S.first() # using first as arbitrary pivot
8     L = LinkedQueue()
9     E = LinkedQueue()
10    G = LinkedQueue()
11    while not S.is_empty(): # divide S into L, E, and G
12        if S.first() < p:
13            L.enqueue(S.dequeue())
14        elif p < S.first():
15            G.enqueue(S.dequeue())
16        else: # S.first() must equal pivot
17            E.enqueue(S.dequeue())
18    # conquer (with recursion)
19    quick_sort(L) # sort elements less than p
20    quick_sort(G) # sort elements greater than p
21    # concatenate results
22    while not L.is_empty():
23        S.enqueue(L.dequeue())
24    while not E.is_empty():
25        S.enqueue(E.dequeue())
26    while not G.is_empty():
27        S.enqueue(G.dequeue())
```


Quick-Sort Tree

- An execution of quick-sort is depicted by a binary tree
 - Each node represents a recursive call of quick-sort and stores
 - Unsorted sequence before the execution and its pivot
 - Sorted sequence at the end of the execution
 - The root is the initial call
 - The leaves are calls on subsequences of size 0 or 1




Execution Example

- Pivot selection




Execution Example (cont.)

- Partition, recursive call, pivot selection




Execution Example (cont.)

- Partition, recursive call, base case




Execution Example (cont.)

- Recursive call, ..., base case, join



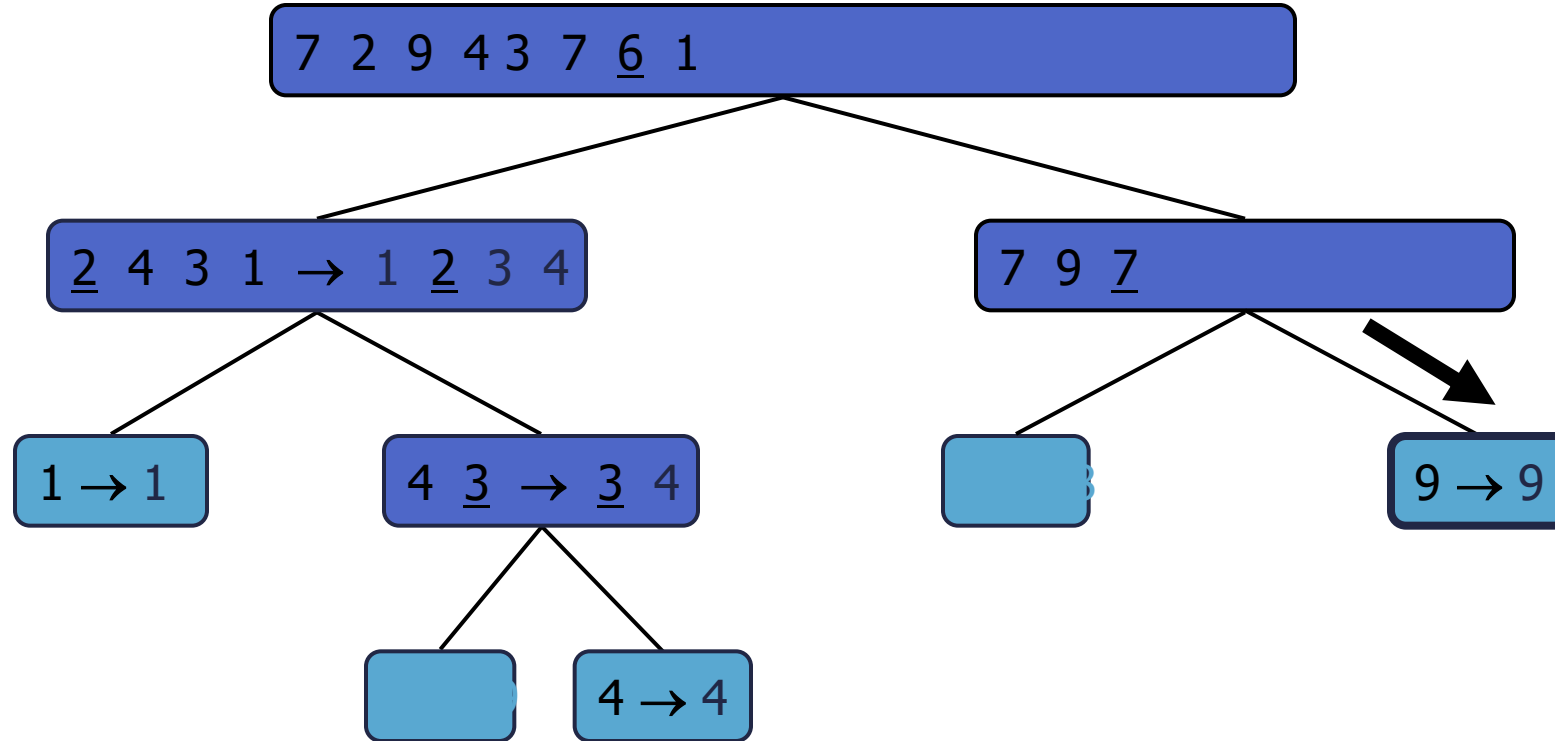
Execution Example (cont.)

- Recursive call, pivot selection




Execution Example (cont.)

- Partition, ..., recursive call, base case



Execution Example (cont.)

- Join, join



Worst-case Running Time

- The worst case for quick-sort occurs when the pivot is the unique minimum or maximum element
- One of L and G has size $n - 1$ and the other has size 0
- The running time is proportional to the sum

$$n + (n - 1) + \dots + 2 + 1$$

- Thus, the worst-case running time of quick-sort is $O(n^2)$


depth time

0 n

1 $n - 1$

...


$n - 1$ 1




Quick-Sort

Expected Running Time

- Consider a recursive call of quick-sort on a sequence of size s
 - Good call: the sizes of L and G are each less than $3s/4$
 - Bad call: one of L and G has size greater than $3s/4$



Good call



Bad call



Expected Running Time, Part 2


- Probabilistic Fact: The expected number of coin tosses required in order to get k heads is $2k$
- For a node of depth i , we expect
 - $i/2$ ancestors are good calls (divide the sequence better than $1/4$ & $3/4$)
 - The size of the input sequence for the current call is at most $(3/4)^{i/2}n$

◆ Therefore, we have

- For a node of depth $2\log_{4/3}n$, the expected input size is one
- The expected height of the quick-sort tree is $O(\log n)$

◆ The amount of work done at the nodes of the same depth is $O(n)$

◆ Thus, the expected running time of quick-sort is $O(n \log n)$



In-Place Quick-Sort




- Quick-sort can be implemented to run in-place
- In the partition step, we use replace operations to rearrange the elements of the input sequence such that
 - the elements less than the pivot have rank less than h
 - the elements equal to the pivot have rank between h and k
 - the elements greater than the pivot have rank greater than k
- The recursive calls consider
 - elements with rank less than h
 - elements with rank greater than k

Algorithm *inPlaceQuickSort*(S, l, r)
Input sequence S , ranks l and r
Output sequence S with the elements of rank between l and r rearranged in increasing order
if $l \geq r$
 return
 $i \leftarrow$ a random integer between l and r
 $x \leftarrow S.\text{elemAtRank}(i)$
 $(h, k) \leftarrow \text{inPlacePartition}(x)$
inPlaceQuickSort($S, l, h - 1$)
inPlaceQuickSort($S, k + 1, r$)


In-Place Partitioning



- Perform the partition using two indices to split S into L and $E \cup G$ (a similar method can split $E \cup G$ into E and G).



- Repeat until j and k cross:
 - Scan j to the right until finding an element $\geq x$.
 - Scan k to the left until finding an element $< x$.
 - Swap elements at indices j and k




Python Implementation

```
1 def inplace_quick_sort(S, a, b):
2     """Sort the list from S[a] to S[b] inclusive using the quick-sort algorithm."""
3     if a >= b: return # range is trivially sorted
4     pivot = S[b] # last element of range is pivot
5     left = a # will scan rightward
6     right = b-1 # will scan leftward
7     while left <= right:
8         # scan until reaching value equal or larger than pivot (or right marker)
9         while left <= right and S[left] < pivot:
10             left += 1
11         # scan until reaching value equal or smaller than pivot (or left marker)
12         while left <= right and pivot < S[right]:
13             right -= 1
14         if left <= right: # scans did not strictly cross
15             S[left], S[right] = S[right], S[left] # swap values
16             left, right = left + 1, right - 1 # shrink range
17
18     # put pivot into its final place (currently marked by left index)
19     S[left], S[b] = S[b], S[left]
20     # make recursive calls
21     inplace_quick_sort(S, a, left - 1)
22     inplace_quick_sort(S, left + 1, b)
```

Summary of Sorting Algorithms

| Algorithm | Time | Notes |
|----------------|---------------------------|--|
| selection-sort | $O(n^2)$ | <ul style="list-style-type: none">▪ in-place▪ slow (good for small inputs) |
| insertion-sort | $O(n^2)$ | <ul style="list-style-type: none">▪ in-place▪ slow (good for small inputs) |
| quick-sort | $O(n \log n)$ expected | <ul style="list-style-type: none">▪ in-place, randomized▪ fastest (good for large inputs) |
| heap-sort | $O(n \log n)$ | <ul style="list-style-type: none">▪ in-place▪ fast (good for large inputs) |
| merge-sort | $O(n \log n)$ | <ul style="list-style-type: none">▪ sequential data access▪ fast (good for huge inputs) |

Bucket-Sort and Radix-Sort





Bucket-Sort

- Let be S be a sequence of n (key, element) items with keys in the range $[0, N - 1]$
 - Bucket-sort uses the keys as indices into an auxiliary array B of sequences (buckets)
 - Phase 1: Empty sequence S by moving each entry (k, o) into its bucket $B[k]$
 - Phase 2: For $i = 0, \dots, N - 1$, move the entries of bucket $B[i]$ to the end of sequence S
 - Analysis:
 - Phase 1 takes $O(n)$ time
 - Phase 2 takes $O(n + N)$ time
- Bucket-sort takes $O(n + N)$ time

Algorithm bucketSort(S):

Input: Sequence S of entries with integer keys in the range $[0, N - 1]$

Output: Sequence S sorted in nondecreasing order of the keys

let B be an array of N sequences, each of which is initially empty

for each entry e in S **do**

k = the key of e

 remove e from S

 insert e at the end of bucket $B[k]$

for $i = 0$ to $N - 1$ **do**

for each entry e in $B[i]$ **do**


 remove e from $B[i]$

 insert e at the end of S


Example




- Key range $[0, 9]$



Phase 1



Phase 2



Properties and Extensions



- Key-type Property

- The keys are used as indices into an array and cannot be arbitrary objects
- No external comparator


- Stable Sort Property

- The relative order of any two items with the same key is preserved after the execution of the algorithm

Extensions

- Integer keys in the range $[a, b]$
 - Put entry (k, o) into bucket $B[k - a]$
- String keys from a set D of possible strings, where D has constant size (e.g., names of the 50 U.S. states)
 - Sort D and compute the rank $r(k)$ of each string k of D in the sorted sequence
 - Put entry (k, o) into bucket $B[r(k)]$

Lexicographic Order



- A d -tuple is a sequence of d keys (k_1, k_2, \dots, k_d) , where key k_i is said to be the i -th dimension of the tuple
- Example:
 - The Cartesian coordinates of a point in space are a 3-tuple
- The lexicographic order of two d -tuples is recursively defined as follows

$$(x_1, x_2, \dots, x_d) < (y_1, y_2, \dots, y_d)$$
$$\Leftrightarrow$$

$$x_1 < y_1 \vee x_1 = y_1 \wedge (x_2, \dots, x_d) < (y_2, \dots, y_d)$$

I.e., the tuples are compared by the first dimension, then by the second dimension, etc.

Lexicographic-Sort

- Let C_i be the comparator that compares two tuples by their i -th dimension
- Let $stableSort(S, C)$ be a stable sorting algorithm that uses comparator C
- Lexicographic-sort sorts a sequence of d -tuples in lexicographic order by executing d times algorithm $stableSort$, one per dimension
- Lexicographic-sort runs in $O(dT(n))$ time, where $T(n)$ is the running time of $stableSort$

Algorithm *lexicographicSort*(S)

Input sequence S of d -tuples

Output sequence S sorted in
lexicographic order

for $i \leftarrow d$ **downto** 1


stableSort(S, C_i)

Example:

(7,4,6) (5,1,5) (2,4,6) (2, 1, 4) (3, 2, 4)
(2, 1, **4**) (3, 2, **4**) (5,1,**5**) (7,4,**6**) (2,4,**6**)
(2, **1**, 4) (5,**1**,5) (3, **2**, 4) (7,**4**,6) (2,**4**,6)
(**2**, 1, 4) (**2**,4,6) (**3**, 2, 4) (**5**,1,5) (**7**,4,6)

Radix-Sort

- Radix-sort is a specialization of lexicographic-sort that uses bucket-sort as the stable sorting algorithm in each dimension
- Radix-sort is applicable to tuples where the keys in each dimension i are integers in the range $[0, N - 1]$
- Radix-sort runs in time $O(d(n + N))$




Algorithm *radixSort*(S, N)

Input sequence S of d -tuples such
that $(0, \dots, 0) \leq (x_1, \dots, x_d)$ and
 $(x_1, \dots, x_d) \leq (N - 1, \dots, N - 1)$
for each tuple (x_1, \dots, x_d) in S

Output sequence S sorted in
lexicographic order

for $i \leftarrow d$ **downto** 1
 bucketSort(S, N)

Radix-Sort for Binary Numbers



- Consider a sequence of n b -bit integers

$$x = x_{b-1} \dots x_1 x_0$$

- We represent each element as a b -tuple of integers in the range $[0, 1]$ and apply radix-sort with $N = 2$
- This application of the radix-sort algorithm runs in $O(bn)$ time
- For example, we can sort a sequence of 32-bit integers in linear time

Algorithm *binaryRadixSort*(S)

Input sequence S of b -bit integers


Output sequence S sorted
replace each element x of S with the item $(0, x)$

for $i \leftarrow 0$ to $b - 1$


replace the key k of each item (k, x) of S with bit x_i of x

bucketSort($S, 2$)

Example



- Sorting a sequence of 4-bit integers



Selection





The Selection Problem

- Given an integer k and n elements x_1, x_2, \dots, x_n , taken from a total order, find the k -th smallest element in this set.
- Of course, we can sort the set in $O(n \log n)$ time and then index the k -th element.


$k=3$


7 4 9 6 2 → 2 4 6 7 9

- Can we solve the selection problem faster?

Quick-Select

- Quick-select is a randomized selection algorithm based on the prune-and-search paradigm:
 - Prune: pick a random element x (called **pivot**) and partition S into
 - L : elements less than x
 - E : elements equal x
 - G : elements greater than x
 - Search: depending on k , either answer is in E , or we need to recur in either L or G





Partition

- We partition an input sequence as in the quick-sort algorithm:
 - We remove, in turn, each element y from S and
 - We insert y into L , E or G , depending on the result of the comparison with the pivot x
- Each insertion and removal is at the beginning or at the end of a sequence, and hence takes $O(1)$ time
- Thus, the partition step of quick-select takes $O(n)$ time

Algorithm *partition*(S, p)

Input sequence S , position p of pivot

Output subsequences L , E , G of the elements of S less than, equal to, or greater than the pivot, resp.

$L, E, G \leftarrow$ empty sequences

$x \leftarrow S.remove(p)$

while $\neg S.isEmpty()$

$y \leftarrow S.remove(S.first())$

if $y < x$

$L.addLast(y)$

else if $y = x$

$E.addLast(y)$


else $\{ y > x \}$

$G.addLast(y)$

return L, E, G


Quick-Select Visualization

- An execution of quick-select can be visualized by a recursion path
 - Each node represents a recursive call of quick-select, and stores k and the remaining sequence




Selection


Expected Running Time




- Consider a recursive call of quick-select on a sequence of size s
 - Good call: the sizes of L and G are each less than $3s/4$
 - Bad call: one of L and G has size greater than $3s/4$




Good call



Bad call



Expected Running Time, Part 2



- Probabilistic Fact #1: The expected number of coin tosses required in order to get one head is two
- Probabilistic Fact #2: Expectation is a linear function:
 - $E(X + Y) = E(X) + E(Y)$
 - $E(cX) = cE(X)$
- Let $T(n)$ denote the expected running time of quick-select.
- By Fact #2,
 - $T(n) \leq T(3n/4) + bn \cdot (\text{expected \# of calls before a good call})$
- By Fact #1,
 - $T(n) \leq T(3n/4) + 2bn$
- That is, $T(n)$ is a geometric series:
 - $T(n) \leq 2bn + 2b(3/4)n + 2b(3/4)^2n + 2b(3/4)^3n + \dots$
- So $T(n)$ is $O(n)$.
- We can solve the selection problem in $O(n)$ expected time.



Deterministic Selection

- We can do selection in $O(n)$ worst-case time.
- Main idea: recursively use the selection algorithm itself to find a good pivot for quick-select:
 - Divide S into $n/5$ sets of 5 each
 - Find a median in each set
 - Recursively find the median of the “baby” medians.

Min size
for L

| | | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|---|
| 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 |
| 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 |
| 3 | 3 | 3 | 3 | 3 | 3 | 3 | 3 | 3 | 3 | 3 |
| 4 | 4 | 4 | 4 | 4 | 4 | 4 | 4 | 4 | 4 | 4 |
| 5 | 5 | 5 | 5 | 5 | 5 | 5 | 5 | 5 | 5 | 5 |

Min size
for G

Selection