

SE274 Data Structure

Lecture 8: Search Trees, Part 4




(textbook: Chapter 11, Red-Black Tree)

May 6, 2020

Instructor: Sunjun Kim


Information&Communication Engineering, DGIST

Recap: Splay tree




Recap: (2,4) Trees

- A (2,4) tree (also called 2-4 tree or 2-3-4 tree) is a multi-way search with the following properties
 - Node-Size Property: every internal node has at most four children
 - Depth Property: all the external nodes have the same depth
- Depending on the number of children, an internal node of a (2,4) tree is called a 2-node, 3-node or 4-node




Recap: (2,4) Trees: Overflow and Split

- We handle an overflow at a 5-node v with a split operation:
 - let $v_1 \dots v_5$ be the children of v and $k_1 \dots k_4$ be the keys of v
 - node v is replaced nodes v' and v''
 - v' is a 3-node with keys $k_1 k_2$ and children $v_1 v_2 v_3$
 - v'' is a 2-node with key k_4 and children $v_4 v_5$
 - key k_3 is inserted into the parent u of v (a new root may be created)
- The overflow may propagate to the parent node u




Recap: (2,4) Trees: Underflow and Fusion

- Deleting an entry from a node v may cause an **underflow**, where node v becomes a 1-node with one child and no keys
- To handle an underflow at node v with parent u , we consider two cases
- **Case 1: the adjacent siblings of v are 2-nodes**
 - **Fusion operation:** we merge v with an adjacent sibling w and move an entry from u to the merged node v'
 - After a fusion, the underflow may propagate to the parent u



Recap: (2,4) Trees: Underflow and Transfer


- To handle an underflow at node v with parent u , we consider two cases
- Case 2: an adjacent sibling w of v is a 3-node or a 4-node
 - Transfer operation:
 1. we move a child of w to v
 2. we move an item from u to v
 3. we move an item from w to u
 - After a transfer, no underflow occurs



Comparison of Map Implementations


| | Search | Insert | Delete | Notes |
|--------------------|------------------------|------------------------|------------------------|--|
| Hash Table | 1 expected | 1 expected | 1 expected | <ul style="list-style-type: none">no ordered map methodssimple to implement |
| Skip List | $\log n$ high prob. | $\log n$ high prob. | $\log n$ high prob. | <ul style="list-style-type: none">randomized insertionsimple to implement |
| AVL and (2,4) Tree | $\log n$ worst-case | $\log n$ worst-case | $\log n$ worst-case | <ul style="list-style-type: none">complex to implement |

Red-Black Trees




From (2,4) to Red-Black Trees

- A red-black tree is a representation of a (2,4) tree by means of a binary tree whose nodes are colored **red** or **black**
- In comparison with its associated (2,4) tree, a red-black tree has
 - same logarithmic time performance
 - simpler implementation with a single node type




Red-Black Trees

- A red-black tree can also be defined as a binary search tree that satisfies the following properties:
 - Root Property: the root is black
 - External Property: every leaf is black
 - Internal Property: the children of a red node are black
 - Depth Property: all the leaves have the same black depth



Quiz


- Is this a valid red-black tree?



- Root Property: the root is black
- External Property: every leaf is black
- Internal Property: the children of a red node are black
- Depth Property: all the leaves have the same black depth

Quiz


- Is this a valid red-black tree?



- Root Property: the root is black
- External Property: every leaf is black
- Internal Property: the children of a red node are black
- Depth Property: all the leaves have the same black depth

Quiz


- Is this a valid red-black tree?



- Root Property: the root is black
- External Property: every leaf is black
- Internal Property: the children of a red node are black
- Depth Property: all the leaves have the same black depth

Quiz


- Is this a valid red-black tree?



- Root Property: the root is black
- External Property: every leaf is black
- Internal Property: the children of a red node are black
- Depth Property: all the leaves have the same black depth

Quiz

- Is this a valid red-black tree?



- Root Property: the root is black
- External Property: every leaf is black
- Internal Property: the children of a red node are black
- Depth Property: all the leaves have the same black depth

Height of a Red-Black Tree


- **Theorem:** A red-black tree storing n items has height $O(\log n)$

Proof:

- The height of a red-black tree is at most twice the height of its associated (2,4) tree, which is $O(\log n)$
- The search algorithm for a binary search tree is the same as that for a binary search tree
- By the above theorem, searching in a red-black tree takes $O(\log n)$ time

Insertion

- To insert (k, o) , we execute the insertion algorithm for binary search trees and color **red** the newly inserted node z unless it is the root
 - We preserve the root, external, and depth properties
 - If the parent v of z is black, we also preserve the internal property and we are done
 - Else (v is red) we have a **double red** (i.e., a violation of the internal property), which requires a reorganization of the tree
- Example where the insertion of 4 causes a double red:




Remedying a Double Red

- Consider a double red with child z and parent v , and let w be the sibling of v


Case 1: w is black

- The double red is an incorrect replacement of a 4-node
- Restructuring:** we change the 4-node replacement




Case 2: w is red

- The double red corresponds to an overflow
- Recoloring:** we perform the equivalent of a split




Restructuring

- A restructuring remedies a child-parent double red when the parent red node has a black sibling
- It is equivalent to restoring the correct replacement of a 4-node
- The internal property is restored and the other properties are preserved




Restructuring (cont.)

- There are four restructuring configurations depending on whether the double red nodes are left or right children

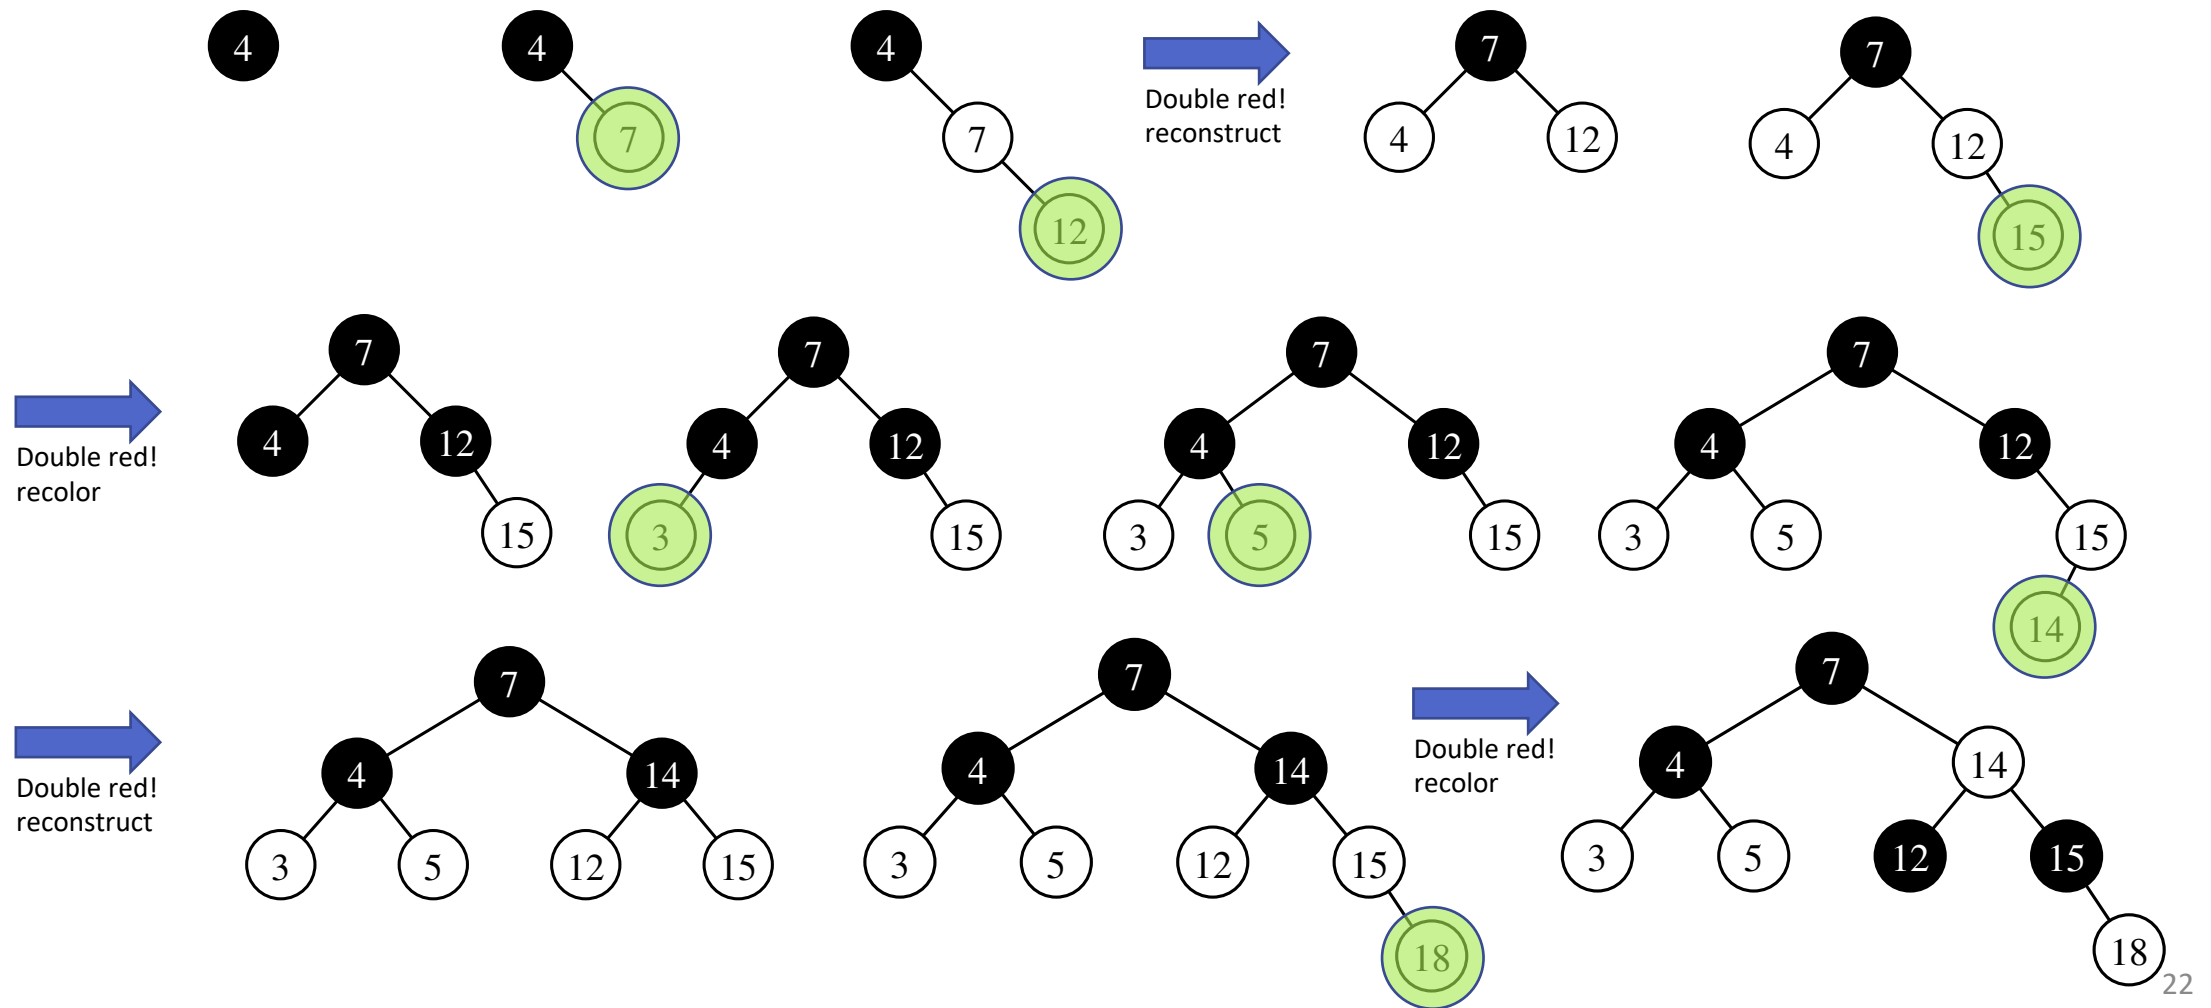


Recoloring

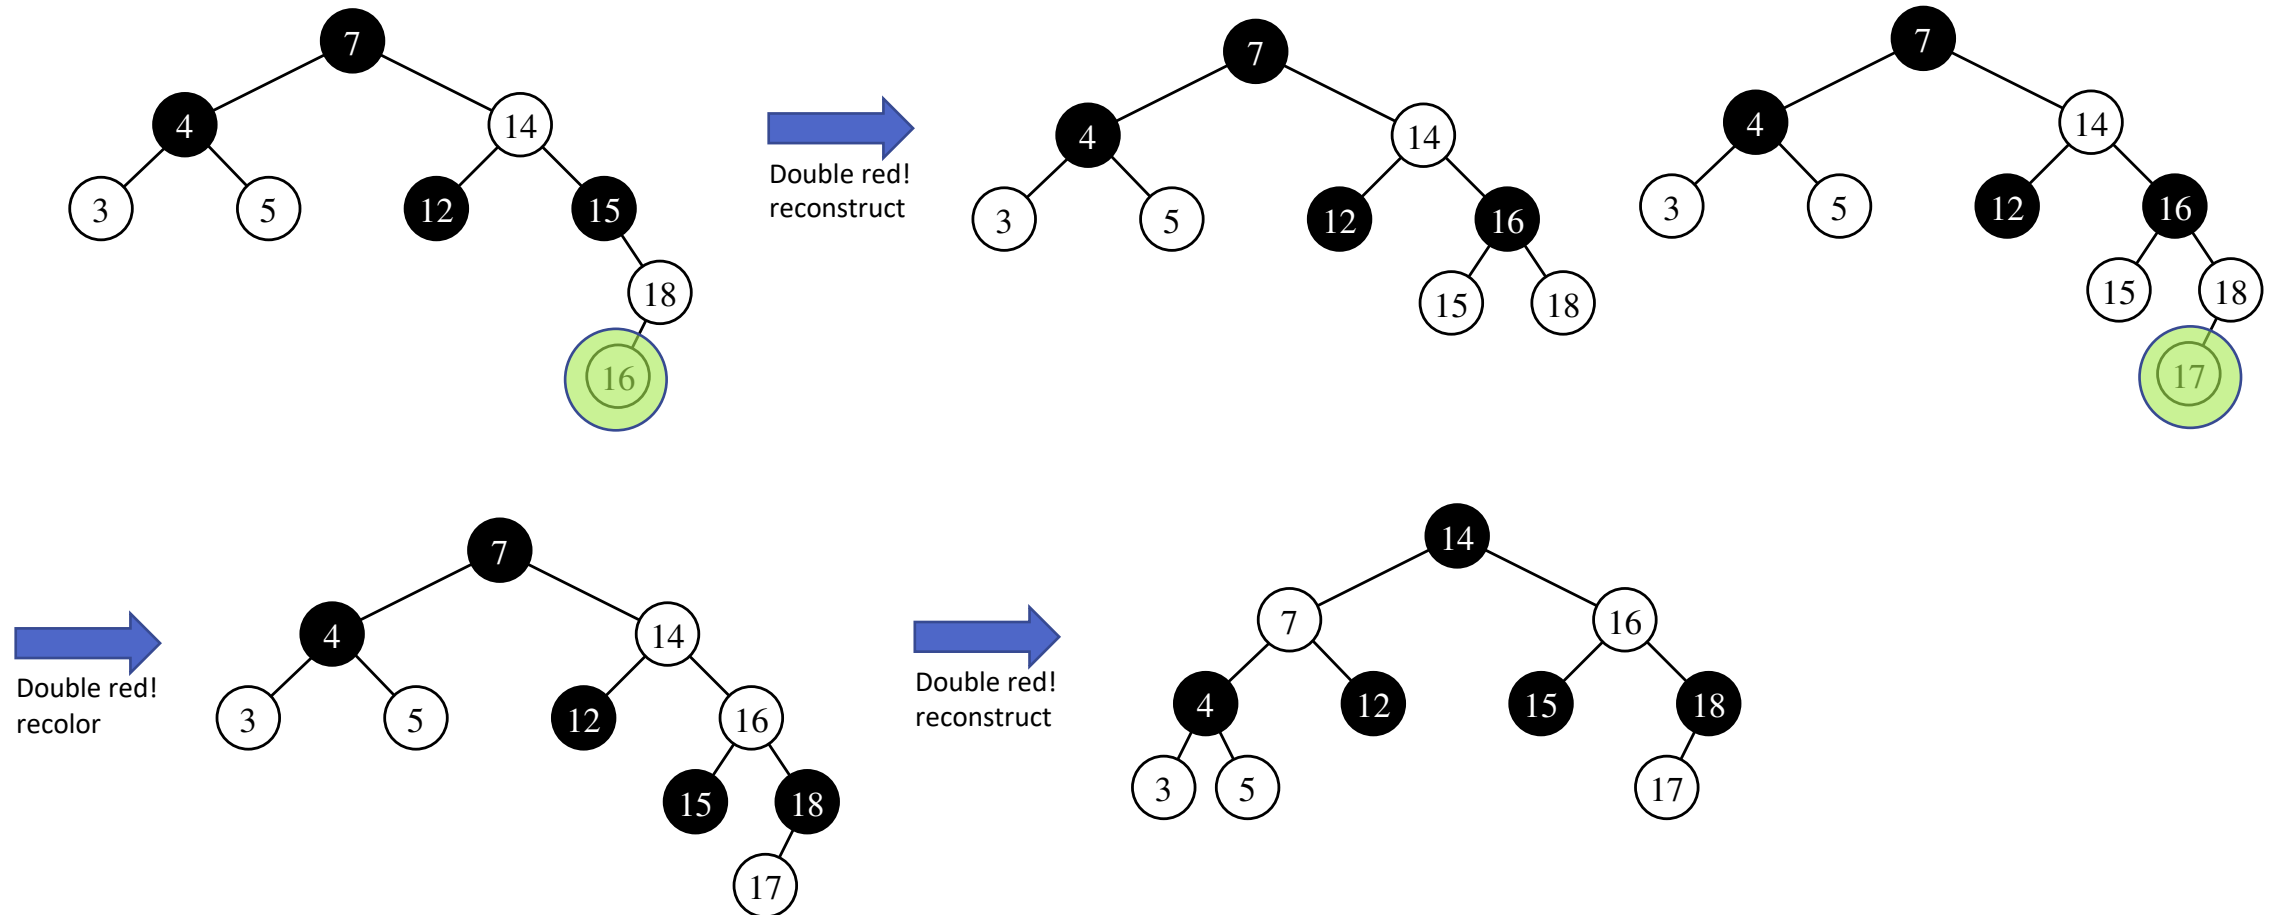
- A recoloring remedies a child-parent double red when the parent red node has a red sibling
- The parent v and its sibling w become black and the grandparent u becomes red, unless it is the root
- It is equivalent to performing a split on a 5-node
- The double red violation may propagate to the grandparent u



Red-Black Insertion: Example



Red-Black Insertion: Example (cont)



Analysis of Insertion

Algorithm *insert(k, o)*


1. We search for key k to locate the insertion node z
2. We add the new entry (k, o) at node z and color z red
3. **while** *doubleRed*(z)
 if *isBlack*(*sibling*(*parent*(z)))
 $z \leftarrow \text{restructure}(z)$
 return
 else { *sibling*(*parent*(z)) is red }
 $z \leftarrow \text{recolor}(z)$

- Recall that a red-black tree has $O(\log n)$ height
- Step 1 takes $O(\log n)$ time because we visit $O(\log n)$ nodes
- Step 2 takes $O(1)$ time
- Step 3 takes $O(\log n)$ time because we perform
 - $O(\log n)$ recolorings, each taking $O(1)$ time, and
 - at most one restructuring taking $O(1)$ time
- Thus, an insertion in a red-black tree takes $O(\log n)$ time

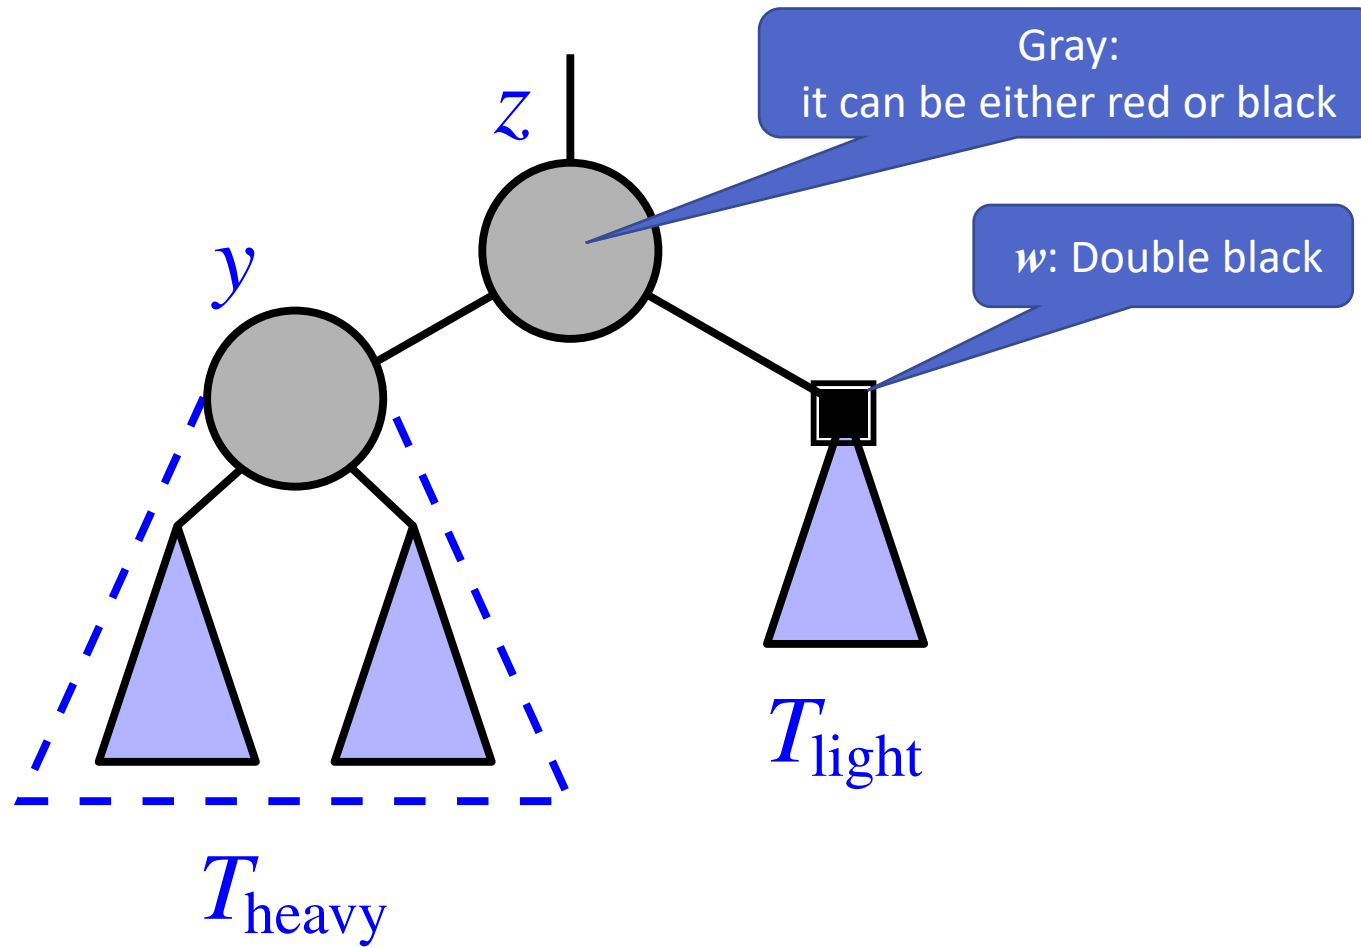
Deletion

- Root Property: the root is black
- External Property: every leaf is black
- Internal Property: the children of a red node are black
- Depth Property: all the leaves have the same black depth


- To perform operation $\text{remove}(k)$, we first execute the deletion algorithm for binary search trees
- Let v be the internal node removed, w the external node removed, and r the sibling of w
 - If either v or r was red, we color r black and we are done
 - Else (v and r were both black) we color r **double black**, which is a violation of the internal property requiring a reorganization of the tree
- Example where the deletion of 8 causes a double black:



T_{heavy} & T_{light}



= has more black depth by 1



Remedying a Double Black

- The algorithm for remedying a double black node w with sibling y considers three cases

Case 1: y is black and has a red child

- We perform a **restructuring**, equivalent to a **transfer**, and we are done

Case 2: y is black and its children are both black

- We perform a **recoloring**, equivalent to a **fusion**, which may propagate up the double black violation

Case 3: y is red


- We perform an **adjustment**, equivalent to choosing a different representation of a 3-node
- After, either Case 1 or Case 2 applies
- Deletion in a red-black tree takes $O(\log n)$ time

Case 1: y is black and has a red child

- The algorithm for remedying a double black node w with sibling y considers three cases

Case 1: y is black and has a red child

- We perform a restructuring, equivalent to a transfer, and we are done




Case 1: y is black and has a red child

- The algorithm for remedying a double black node w with sibling y considers three cases

Case 1: y is black and has a red child

- We perform a restructuring, equivalent to a transfer, and we are done

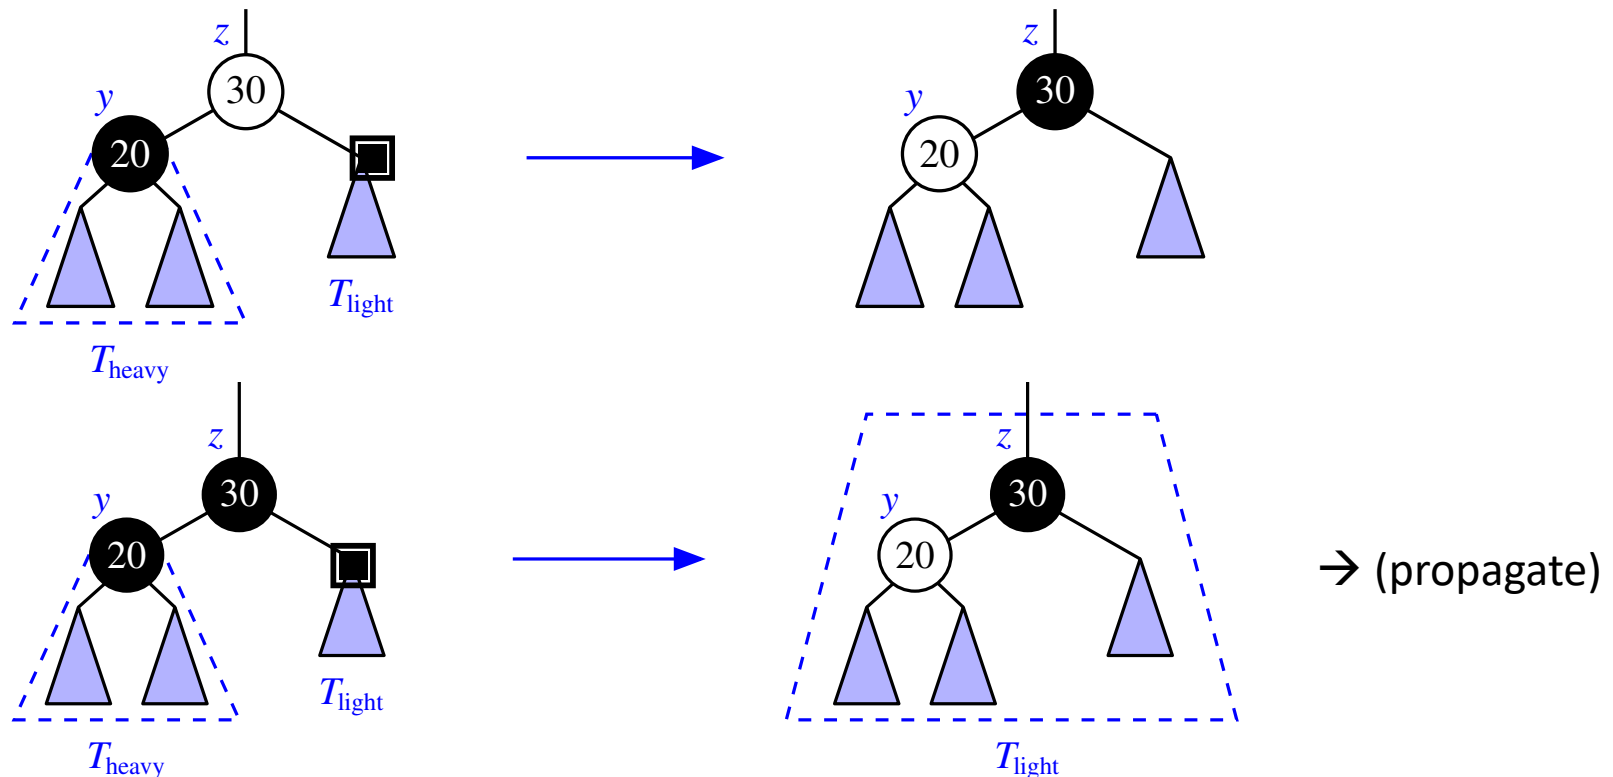


Case 2: y is black and its children are both black

- The algorithm for remedying a double black node w with sibling y considers three cases

Case 2: y is black and its children are both black

- We perform a recoloring, equivalent to a fusion, which may propagate up the double black violation




Case 2: y is black and its children are both black


- The algorithm for remedying a double black node w with sibling y considers three cases

Case 3: y is red

- We perform an **adjustment**, equivalent to choosing a different representation of a 3-node.
- After, either Case 1 or Case 2 applies

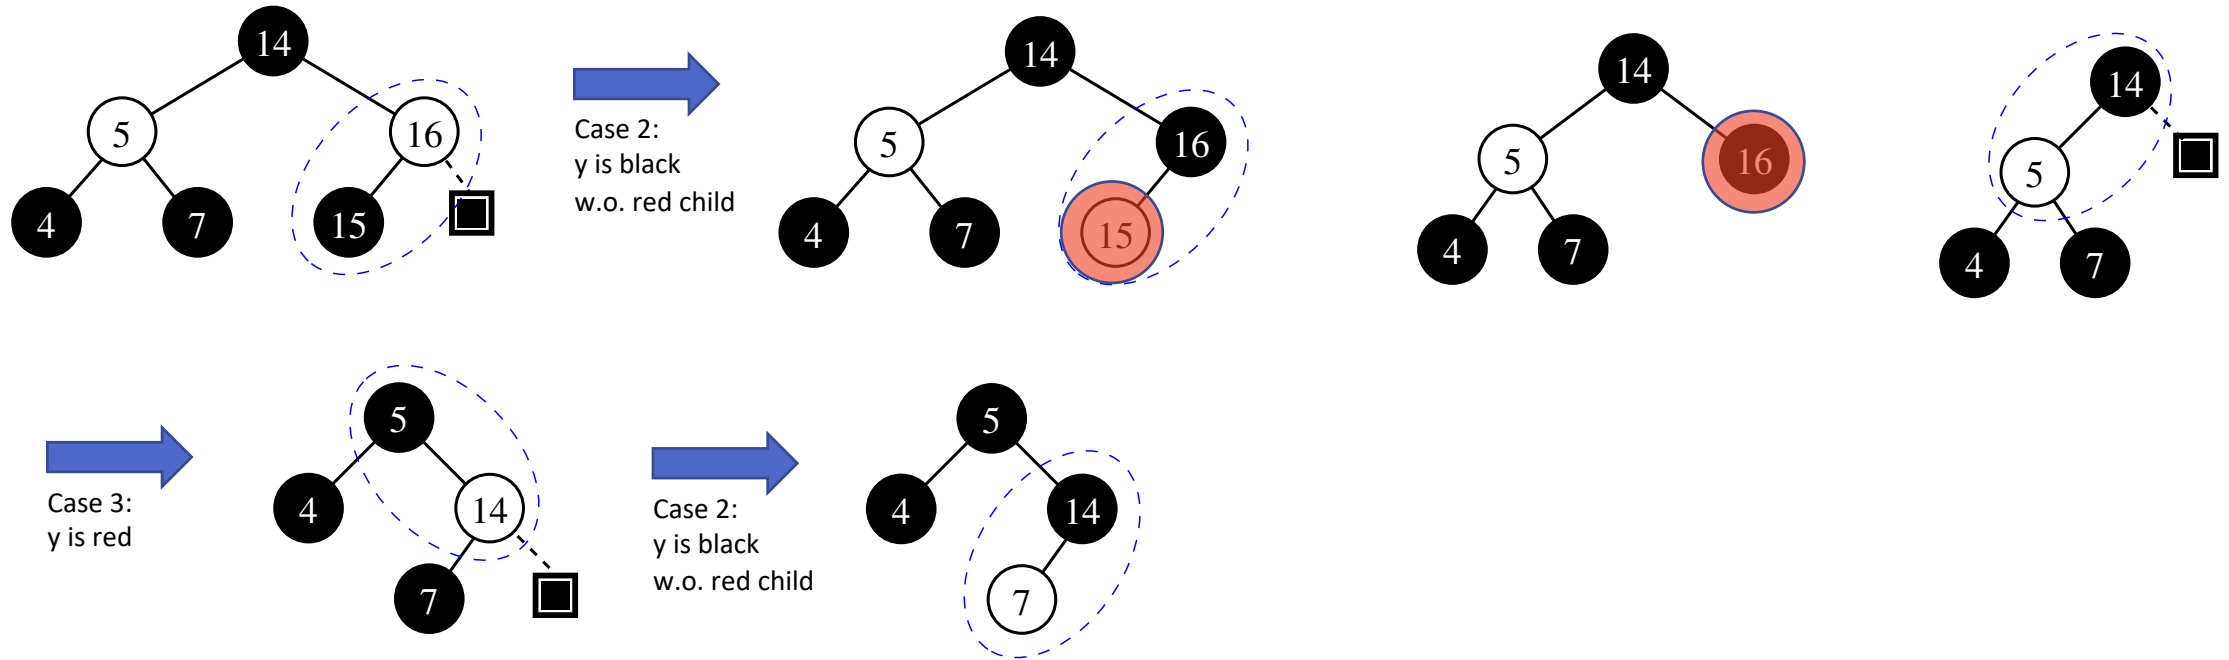


Deletion example



- Case 1: y is black and has a red child \rightarrow reconstruct
- Case 2: y is black and its children are both black \rightarrow recolor
- Case 3: y is red \rightarrow arrangement \rightarrow Case 1 or 2

Deletion example (cont)



- Case 1: y is black and has a red child \rightarrow reconstruct
- Case 2: y is black and its children are both black \rightarrow recolor
- Case 3: y is red \rightarrow arrangement \rightarrow Case 1 or 2

Red-Black Tree Reorganization

| Insertion remedy double red | | |
|------------------------------------|---------------------------------|-------------------------------------|
| Red-black tree action | (2,4) tree action | result |
| restructuring | change of 4-node representation | double red removed |
| recoloring | split | double red removed or propagated up |

| Deletion remedy double black | | |
|-------------------------------------|---------------------------------|---------------------------------------|
| Red-black tree action | (2,4) tree action | result |
| restructuring | transfer | double black removed |
| recoloring | fusion | double black removed or propagated up |
| adjustment | change of 3-node representation | restructuring or recoloring follows |

Python Implementation

```
1 class RedBlackTreeMap(TreeMap):
2     """Sorted map implementation using a red-black tree."""
3     class _Node(TreeMap._Node):
4         """Node class for red-black tree maintains bit that denotes color."""
5         __slots__ = '_red'      # add additional data member to the Node class
6
7     def __init__(self, element, parent=None, left=None, right=None):
8         super().__init__(element, parent, left, right)
9         self._red = True      # new node red by default
```

Python Implementation, Part 2

```
10 #----- positional-based utility methods -----
11 # we consider a nonexistent child to be trivially black
12 def _set_red(self, p): p._node._red = True
13 def _set_black(self, p): p._node._red = False
14 def _set_color(self, p, make_red): p._node._red = make_red
15 def _is_red(self, p): return p is not None and p._node._red
16 def _is_red_leaf(self, p): return self._is_red(p) and self.is_leaf(p)
17
18 def _get_red_child(self, p):
19     """Return a red child of p (or None if no such child)."""
20     for child in (self.left(p), self.right(p)):
21         if self._is_red(child):
22             return child
23     return None
24
25 #----- support for insertions -----
26 def _rebalance_insert(self, p):
27     self._resolve_red(p) # new node is always red
28
29 def _resolve_red(self, p):
30     if self.is_root(p):
31         self._set_black(p) # make root black
32     else:
33         parent = self.parent(p)
34         if self._is_red(parent): # double red problem
35             uncle = self.sibling(parent)
36             if not self._is_red(uncle): # Case 1: misshapen 4-node
37                 middle = self._restructure(p) # do trinode restructuring
38                 self._set_black(middle) # and then fix colors
39                 self._set_red(self.left(middle))
40                 self._set_red(self.right(middle))
41             else: # Case 2: overfull 5-node
42                 grand = self.parent(parent)
43                 self._set_red(grand) # grandparent becomes red
44                 self._set_black(self.left(grand)) # its children become black
45                 self._set_black(self.right(grand))
46                 self._resolve_red(grand) # recur at red grandparent
```

Python Implementation, end

```
47 #----- support for deletions -----
48 def _rebalance_delete(self, p):
49     if len(self) == 1:
50         self._set_black(self.root()) # special case: ensure that root is black
51     elif p is not None:
52         n = self.num_children(p)
53         if n == 1: # deficit exists unless child is a red leaf
54             c = next(self.children(p))
55             if not self._is_red_leaf(c):
56                 self._fix_deficit(p, c)
57         elif n == 2: # removed black node with red child
58             if self._is_red_leaf(self.left(p)):
59                 self._set_black(self.left(p))
60             else:
61                 self._set_black(self.right(p))
62
63 def _fix_deficit(self, z, y):
64     """Resolve black deficit at z, where y is the root of z's heavier subtree."""
65     if not self._is_red(y): # y is black; will apply Case 1 or 2
66         x = self._get_red_child(y)
67         if x is not None: # Case 1: y is black and has red child x; do "transfer"
68             old_color = self._is_red(z)
69             middle = self._restructure(x)
70             self._set_color(middle, old_color) # middle gets old color of z
71             self._set_black(self.left(middle)) # children become black
72             self._set_black(self.right(middle))
73         else: # Case 2: y is black, but no red children; recolor as "fusion"
74             self._set_red(y)
75             if self._is_red(z):
76                 self._set_black(z) # this resolves the problem
77             elif not self.is_root(z):
78                 self._fix_deficit(self.parent(z), self.sibling(z)) # recur upward
79     else: # Case 3: y is red; rotate misaligned 3-node and repeat
80         self._rotate(y)
81         self._set_black(y)
82         self._set_red(z)
83         if z == self.right(y):
84             self._fix_deficit(z, self.left(z))
85         else:
86             self._fix_deficit(z, self.right(z))
```