

SE274 Data Structure

Lecture 11: Memory Management

(textbook: Chapter 15)

June 1, 2020

Instructor: Sunjun Kim

Information&Communication Engineering, DGIST


B-Trees





Computer Memory

- ❑ In order to implement any data structure on an actual computer, we need to use computer memory.
- ❑ Computer memory is organized into a sequence of words, each of which typically consists of 4, 8, or 16 bytes (depending on the computer).
- ❑ These memory words are numbered from 0 to $N - 1$, where N is the number of memory words available to the computer.
- ❑ The number associated with each memory word is known as its memory **address**.



Disk Blocks

- Consider the problem of maintaining a large collection of items that does not fit in main memory, such as a typical database.
- In this context, we refer to the external memory is divided into blocks, which we call **disk blocks**.
- The transfer of a block between external memory and primary memory is a **disk transfer** or **I/O**.
- There is a great time difference that exists between main memory accesses and disk accesses
- Thus, we want to minimize the number of disk transfers needed to perform a query or update. We refer to this count as the **I/O complexity** of the algorithm involved.

(a,b) Trees

- To reduce the number of external-memory accesses when searching, we can represent a map using a multiway search tree.
- This approach gives rise to a generalization of the (2,4) tree data structure known as the **(a,b) tree**.
- An (a,b) tree is a multiway search tree such that each node has between a and b children and stores between $a - 1$ and $b - 1$ entries.
- By setting the parameters a and b appropriately with respect to the size of disk blocks, we can derive a data structure that achieves good external-memory performance.

Definition

- An **(a,b) tree**, where parameters a and b are integers such that $2 \leq a \leq (b+1)/2$, is a multiway search tree T with the following additional restrictions:
- **Size Property**: Each internal node has at least a children, unless it is the root, and has at most b children.
- **Depth Property**: All the external nodes have the same depth.

Height of an (a,b) Tree

Proposition 15.1: *The height of an (a,b) tree storing n entries is $\Omega(\log n / \log b)$ and $O(\log n / \log a)$.*

Justification: Let T be an (a,b) tree storing n entries, and let h be the height of T . We justify the proposition by establishing the following bounds on h :

$$\frac{1}{\log b} \log(n+1) \leq h \leq \frac{1}{\log a} \log \frac{n+1}{2} + 1.$$

By the size and depth properties, the number n'' of external nodes of T is at least $2a^{h-1}$ and at most b^h . By Proposition 11.7, $n'' = n+1$. Thus,

$$2a^{h-1} \leq n+1 \leq b^h.$$

Taking the logarithm in base 2 of each term, we get

$$(h-1) \log a + 1 \leq \log(n+1) \leq h \log b.$$

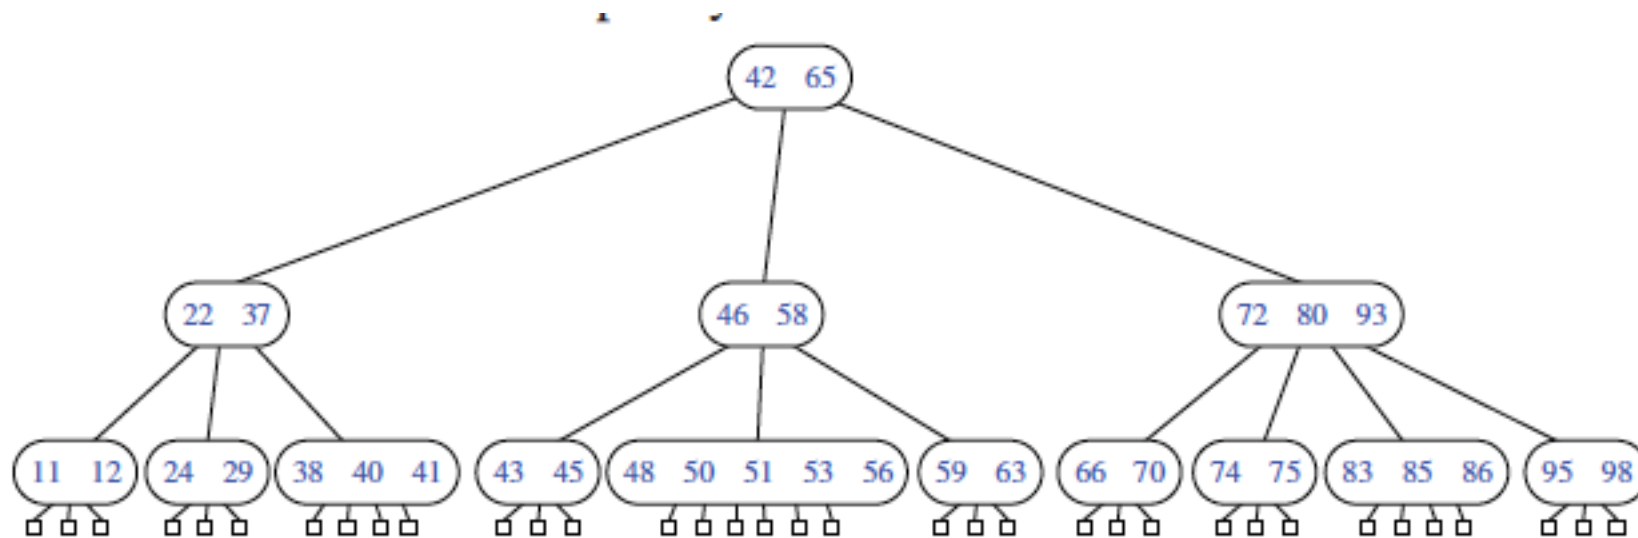
An algebraic manipulation of these inequalities completes the justification. ■

Searches and Updates

- The search algorithm in an **(a,b)** tree is exactly like the one for multiway search trees.
- The insertion algorithm for an **(a,b)** tree is similar to that for a **(2,4)** tree.
 - An overflow occurs when an entry is inserted into a **b**-node **w**, which becomes an illegal **(b+1)**-node.
 - To remedy an overflow, we split node **w** by moving the median entry of **w** into the parent of **w** and replacing **w** with a **(b+1)/2**-node **w** and a **(b+1)/2**-node **w**.
- Removing an entry from an **(a,b)** tree is similar to what was done for **(2,4)** trees.
 - An underflow occurs when a key is removed from an **a**-node **w**, distinct from the root, which causes **w** to become an **(a-1)**-node.
 - To remedy an underflow, we perform a transfer with a sibling of **w** that is not an **a**-node or we perform a fusion of **w** with a sibling that is an **a**-node.

B-Trees

- A version of the **(a,b)** tree data structure, which is the best-known method for maintaining a map in external memory, is a “**B-tree**.”
- A **B-tree of order d** is an **(a,b)** tree with $a = d/2$ and $b = d$.



I/O Complexity

Proposition 15.2: *A B-tree with n entries has I/O complexity $O(\log_B n)$ for search or update operation, and uses $O(n/B)$ blocks, where B is the size of a block.*

- **Proof:**
 - Each time we access a node to perform a search or an update operation, we need only perform a single disk transfer.
 - Each search or update requires that we examine at most **$O(1)$** nodes for each level of the tree.

Data Structure Final Exam

Final Exam (재공지)

- Date: 6월 9일 (화요일)
- Time: 10:00 – 11:30 (90 min)
- Location: E1 컨벤션 A, 컨벤션 B
- Details:
 - 시험범위: 9 주차 (Search Tree) – 15주차 내용
 - Open-book exam
 - 전자기기를 제외한 모든 Material 허용
 - 책 전체를 가져오기보다, 요점을 정리한 cheat-sheet 작성을 추천
 - 답안은 수기 작성 (정자로 알아보기 쉽도록.)